

Frontalization of Profile Face Images Using a Generative Adversarial Network

Abhinav Sinha*
230042
abhinavsi23

Ashish Kumar Soni†
230228
sashish23

Rudransh Verma‡
230881
rudranshv23

Tanvi Pooranmal Meena§
231085
mtanvi23

Course Code & Name: EE655: Computer Vision & Deep Learning
Course Instructor: Prof. Koteswar Rao Jerripothula

Abstract

Face frontalization—the task of synthesizing a front-facing image from a profile view—remains challenging under extreme pose variations due to occlusions and the loss of geometric information. Prior approaches often rely on explicit 3D modeling or landmark detection, which degrade under real-world conditions. In this work, we propose a purely learning-based approach that frames face frontalization as an image-to-image translation problem using a conditional Generative Adversarial Network (cGAN). We construct a high-quality dataset of profile-frontal pairs from the 300W-LP dataset, including flipped variants for pose diversity. Our generator is based on a U-Net architecture that integrates both adversarial and pixel-wise reconstruction losses to ensure realism and identity preservation. Experiments demonstrate that our model can successfully reconstruct plausible frontal views even from challenging side profiles, without requiring 3D priors or identity labels. This provides a lightweight and effective solution for downstream tasks like face recognition and alignment in unconstrained environments.

1 Introduction

Face frontalization is a key task in computer vision with applications in face recognition, alignment, and expression analysis. However, real-world conditions such as pose variation and occlusion introduce significant challenges. Traditional methods often rely on 3D modeling, which can be computationally expensive and brittle under occlusion. In this work, we adopt a fully learning-based approach that avoids the need for any 3D priors or landmarks. By leveraging the power of conditional GANs and the 300W-LP dataset, we aim to generate high-fidelity frontal views from side profiles using end-to-end training. Our approach is simple yet robust, and we demonstrate that it produces realistic and identity-preserving results.

- **3D Morphable Models (3DMM):** Methods using 3DMM for explicit shape/texture fitting and frontal view synthesis [?, ?]. These require accurate landmark alignment and model parameter optimization.
- **Landmark-based Warping:** Approaches leveraging facial landmarks for feature alignment [?, ?]. These methods face challenges with occlusions and resolution limitations due to dependency on landmark accuracy.
- **GAN-based Methods:** State-of-the-art techniques using adversarial networks for frontal view synthesis [?, ?, ?, ?]. DA-GAN [?] introduces dual attention mechanisms, while FFWM [?] combines flow-based warping with GANs.

Our approach builds upon the GAN-based paradigm, but without any auxiliary labels or pose annotations. This simplifies training while maintaining quality.

* Department of Mechanical Engineering

† Department of Mathematics & Statistics

‡ Department of Computer Science & Engineering

§ Department of Mechanical Engineering

2 Proposed Method

We formulate face frontalization as a conditional image-to-image translation problem, where the model learns a mapping from profile face images to their corresponding frontal views. Our approach uses a conditional Generative Adversarial Network (cGAN) framework, comprising a generator G and a discriminator D , trained in an adversarial setting.

2.1 Generator Architecture

The generator G is modeled as a U-Net [?], which is particularly suited for image-to-image translation tasks due to its encoder-decoder structure with skip connections. The encoder progressively downsamples the input image to a low-dimensional latent representation, while the decoder upsamples it back to the original resolution. Skip connections between corresponding layers in the encoder and decoder help preserve spatial features and fine details.

- **Input:** $256 \times 256 \times 3$ RGB profile face image.
- **Encoder:** 8 convolutional layers with increasing filter sizes $\{64, 128, 256, 512, 512, 512, 512\}$, each followed by LeakyReLU and InstanceNorm.
- **Decoder:** 8 deconvolutional layers with symmetric filter sizes, using ReLU, InstanceNorm, and Dropout (in early layers).
- **Output:** $256 \times 256 \times 3$ RGB image representing the frontalized face.

2.2 Discriminator Architecture

The discriminator D is designed as a PatchGAN, which classifies overlapping patches of the image rather than the entire image. This encourages high-frequency detail preservation and local realism.

- **Input:** Concatenated pair of input profile and real/generated frontal image.
- **Architecture:** 5 convolutional layers with increasing filters $\{64, 128, 256, 512, 1\}$, each followed by LeakyReLU and InstanceNorm (except final layer).
- **Output:** $N \times N$ probability map where each value indicates whether the corresponding patch is real or fake.

2.3 Loss Functions

The model is trained with a weighted combination of adversarial loss and L1 reconstruction loss:

$$\mathcal{L} = \mathcal{L}_{adv}(G, D) + \lambda_1 \mathcal{L}_{L1}(G)$$

Adversarial Loss

The adversarial loss encourages the generator to produce realistic frontal faces:

$$\mathcal{L}_{adv}(G, D) = \mathbb{E}_{x,y}[\log D(x, y)] + \mathbb{E}_x[\log(1 - D(x, G(x)))]$$

where x is the input profile face and y is the corresponding ground truth frontal face.

L1 Loss

The L1 loss ensures that the generated image is close to the ground truth at a pixel level, preserving structure and identity:

$$\mathcal{L}_{L1}(G) = \mathbb{E}_{x,y}[\|y - G(x)\|_1]$$

We empirically set $\lambda_1 = 100$ to balance realism with pixel-level accuracy, following common practices in image-to-image translation tasks.

2.4 Training details

The following hyperparameters were used in the model:

- **Optimizer:** Adam
- **Learning rate:** $2e-4$
- **Beta 1:** 0.5
- **Loss function:** Combined (Adversarial + Pixel-wise)
- **Adversarial loss scaling:** 1 (implicitly)
- **Pixel-wise loss scaling:** 100
- **Input Image Size:** 128×128
- **Number of epochs:** 10
- **Batch size:** 32
- **Generator Architecture:** U-Net
- **Number of filters (Conv2D/Conv2DTranspose):** 64, 128, 256, 512
- **Kernel size (Conv2D/Conv2DTranspose):** 4×4
- **Strides (Conv2D/Conv2DTranspose):** 2
- **Discriminator Architecture:** PatchGAN (similar to U-Net)
- **Number of filters (Conv2D):** 64, 128, 256, 512
- **Strides (Conv2D):** 2

- **Data Preprocessing:** Resize to 128x128, Normalize to $[-1, 1]$

The model is implemented in PyTorch and trained using mixed precision to reduce memory usage and speed up convergence. Training is done on an NVIDIA RTX GPU with 12GB VRAM, and training time per epoch is approximately 4 minutes.

3 Dataset and Preprocessing

We use the 300W-LP dataset, which is a large-scale dataset consisting of profile and frontal face pairs generated through pose synthesis from the Multi-PIE dataset. This dataset provides a diverse range of facial images across different poses, which is crucial for training a robust face frontalization model. We begin by downloading the dataset from the official repository, which contains multiple pose variations for each individual face.

In the dataset, each subject is represented by 24 images: one frontal image and 23 images corresponding to different side poses. These side poses include variations in yaw, pitch, and roll, covering a wide range of facial orientations. For each individual, we create pairs consisting of one frontal image and the corresponding profile images from different side views. Specifically, we ensure that the data includes images taken from both left and right side views by flipping the profile images horizontally. This augmentation helps to further increase the pose diversity, making the model more robust to different orientations.

The dataset consists of a total of 114,000 images after pairing the frontal and profile views for all individuals.

We then preprocess the images to ensure they are in a consistent format suitable for training:

- **Resize:** All images are resized to a resolution of 256×256 to maintain consistency and allow for efficient training. This resolution strikes a balance between computational efficiency and the ability to capture facial details.
- **Normalization:** The pixel values of the images are normalized to the range $[-1, 1]$ to improve the training stability and facilitate the convergence of the neural network.
- **Data Augmentation:** To further enrich the dataset and prevent overfitting, we apply random horizontal flipping to profile images, which simulates mirror poses. This allows the model to learn to generate frontal faces from both left and right orientations.

We then divide the dataset into training and testing subsets. Specifically, 80% of the images are used for training, and the remaining 20% are held out for testing and evaluation. The training set includes a diverse range of profiles and their corresponding frontal images, while the test set is used to evaluate the model’s ability to generalize to new, unseen profiles.

In summary, the dataset is processed as follows:

1. Download and extract the 300W-LP dataset, containing frontal and profile image pairs.
2. Pair one frontal image with 23 different side poses per individual, totaling 24 images per face.
3. Augment the data by flipping profile images horizontally to simulate both left and right side views.
4. Resize all images to 256×256 and normalize the pixel values to the range $[-1, 1]$.
5. Split the dataset into training (80%) and testing (20%) subsets.

This preprocessing pipeline ensures that the model is trained on a diverse and well-prepared dataset, enabling it to generate high-quality frontal faces from profile images even under challenging conditions such as extreme pose variations and occlusion.

4 Experiments and Results

We trained our model on an NVIDIA GPU for 10 epochs using the Adam optimizer.

Key hyperparameters include a learning rate of 2×10^{-4} and $\lambda_1 = 100$ for the L1 loss.

To evaluate the performance of the model, we performed both qualitative and quantitative assessments.

4.1 Qualitative Evaluation

The results are presented through side-by-side comparisons of input profile images and generated frontal images. These visualizations provide an insight into the model’s ability to effectively reconstruct the frontal view from profile images.

4.2 Quantitative Evaluation

We evaluate the frontalization quality using standard image quality metrics, such as Peak Signal-to-Noise Ratio (PSNR) and Structural Similarity Index (SSIM). The following table shows the quantitative evaluation results:

Table 1: Quantitative evaluation of frontalization quality using PSNR and SSIM.

Metric	Score
PSNR	22.1 dB
SSIM	0.65

4.3 Discussion

Based on the qualitative and quantitative results, the model demonstrates a good balance between high-quality frontalization and reasonable computational efficiency. The results suggest that the model performs effectively in reconstructing frontal views from profile images.

5 Application and Novelty

The proposed method has several potential applications in the field of computer vision, particularly in the domain of facial recognition, face synthesis, and augmented reality. By generating frontal face images from profile images, this approach can be used to enhance facial recognition systems, enabling them to accurately recognize faces from varying orientations. This has applications in surveillance, security, and authentication systems where a subject’s profile view is the only available image.

Additionally, the technique can be applied in virtual environments, such as in video games and augmented reality, where realistic 3D models of faces can be generated from a limited set of 2D profile images. This opens up new possibilities for user interactions, character modeling, and virtual facial expressions. The model could also be used for applications in digital forensics, where frontal face generation could aid in the identification of individuals from partial or distorted images.

5.1 Novelty

The novelty of our approach lies in several aspects:

- **Pose-Invariant Frontalization:** We frame face frontalization as an image-to-image translation task, utilizing a U-Net architecture that includes skip connections to preserve spatial information. The model is capable of generating high-quality frontal face images from profile images, even in the case of extreme yaw, pitch, and roll variations.
- **Efficient Discriminator:** The use of a PatchGAN as the discriminator is a key contribution. Instead of classifying the entire image as real or fake, PatchGAN classifies each patch of the image,

enabling finer-grained learning. This not only improves the quality of the generated images but also accelerates the training process.

- **Custom Dataset and Augmentation:** We use the 300W-LP dataset, which includes a variety of facial poses derived from the Multi-PIE dataset. By flipping profile images horizontally, we introduce additional pose variation, further improving the robustness of the model. The dataset consists of 114,000 images, and the resulting model is trained on a diverse set of profiles and corresponding frontal views.
- **End-to-End Training with Combined Loss Function:** The training procedure employs a combined loss function that integrates adversarial loss with pixel-wise reconstruction loss, ensuring that the generator not only produces realistic images but also accurately reconstructs the frontal face from the profile. This dual loss formulation is essential for achieving high-quality and realistic outputs.
- **Scalable and Parallelizable Architecture:** The model is implemented using TensorFlow and Keras, and the training process is optimized for both CPU and GPU environments. We leverage the capabilities of TensorFlow for efficient batching and data prefetching, ensuring fast and scalable training.

6 Conclusion

We presented a GAN-based approach for face frontalization that requires no 3D priors or identity labels. By formulating the task as an image-to-image translation problem, we train a simple yet effective cGAN to generate identity-preserving frontal views. Our method performs robustly under extreme pose variation and serves as a practical solution for real-world applications.

Appendix

Code Repository

The complete implementation of our model and training scripts is available at: <https://github.com/Face-Frontalization/Face-Frontalization>

Dataset

300W-LP dataset:

<http://www.cbsr.ia.ac.cn/users/xiangyuzhu/projects/3DDFA/Database/300W-LP/main.htm>

References

- [1] Blanz, V., and Vetter, T. (1999). A morphable model for the synthesis of 3D faces. In Proceedings of SIGGRAPH.
- [2] Yin, X., Yu, X., Sohn, K., Liu, X., Chandraker, M. (2017). Towards large-pose face frontalization in the wild. In Proceedings of the IEEE International Conference on Computer Vision (ICCV).
- [3] Hassner, T., Harel, S., Paz, E., and Enbar, R. (2015). Effective face frontalization in unconstrained images. In Proceedings of CVPR.
- [4] Zhu, X., Lei, Z., Liu, X., Shi, H., and Li, S.Z. (2015). Face alignment across large poses: A 3D solution. In Proceedings of CVPR.
- [5] Tran, L., Yin, X., and Liu, X. (2017). Disentangled representation learning GAN for pose-invariant face recognition. In Proceedings of CVPR.
- [6] Huang, R., Zhang, S., Li, T., and He, R. (2017). Beyond face rotation: Global and local perception GAN for photorealistic and identity preserving frontal view synthesis. In Proceedings of ICCV.